

Neural Networks Deep Learning – Icp6

Name: Deekshitha, Gaddameedhi

ID: 700755765

GitHub link: https://github.com/deekshitha430/icp7_neural

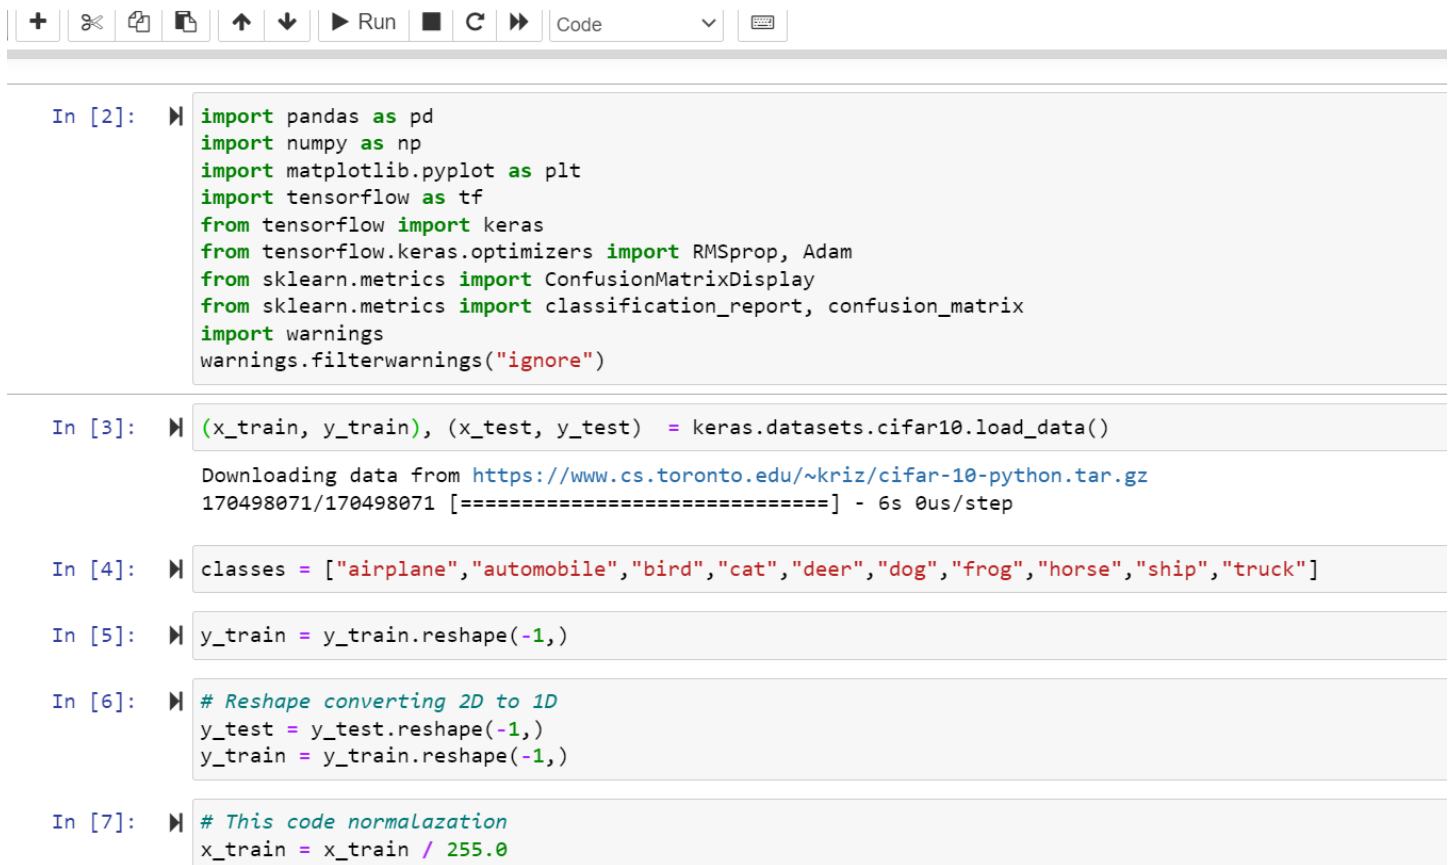
Video Link: https://drive.google.com/file/d/1I7IP-OCO_VdnmXTnYWPMInPcxutW2769/view?usp=sharing

Use Case Description:

LeNet5, AlexNet, Vgg16, Vgg19

1. Training the model
2. Evaluating the model

LeNet5&AlexNet:



```
+  ⏮  ⏭  ⏪  ⏩  ⏴  ⏵  ▶ Run  ⏹  ↺  ⏴▶ Code  ⌵  📄

In [2]: ▶ import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.optimizers import RMSprop, Adam
from sklearn.metrics import ConfusionMatrixDisplay
from sklearn.metrics import classification_report, confusion_matrix
import warnings
warnings.filterwarnings("ignore")

In [3]: ▶ (x_train, y_train), (x_test, y_test) = keras.datasets.cifar10.load_data()

Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170498071/170498071 [=====] - 6s 0us/step

In [4]: ▶ classes = ["airplane", "automobile", "bird", "cat", "deer", "dog", "frog", "horse", "ship", "truck"]

In [5]: ▶ y_train = y_train.reshape(-1,)

In [6]: ▶ # Reshape converting 2D to 1D
y_test = y_test.reshape(-1,)
y_train = y_train.reshape(-1,)

In [7]: ▶ # This code normalization
x_train = x_train / 255.0
```

Neural Networks Deep Learning – Icp6

```
Code
```

```
x_train = x_train / 255.0  
x_test = x_test / 255.0
```

```
[8]: x_train.shape
```

```
Out[8]: (50000, 32, 32, 3)
```

```
[9]: from tensorflow.keras import layers, models  
lenet = keras.models.Sequential([  
    keras.layers.Conv2D(6, kernel_size=5, strides=1, activation='relu', input_shape=(32,32,3), padding='same'), #C1  
    keras.layers.AveragePooling2D(), #S1  
    keras.layers.Conv2D(16, kernel_size=5, strides=1, activation='relu', padding='valid'), #C2  
    keras.layers.AveragePooling2D(), #S2  
    keras.layers.Conv2D(120, kernel_size=5, strides=1, activation='relu', padding='valid'), #C3  
    keras.layers.Flatten(), #Flatten  
    keras.layers.Dense(84, activation='relu'), #F1  
    keras.layers.Dense(10, activation='softmax') #Output layer  
)
```

```
[10]: lenet.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 6)	456
average_pooling2d (AveragePooling2D)	(None, 16, 16, 6)	0
conv2d_1 (Conv2D)	(None, 12, 12, 16)	2416
average_pooling2d_1 (AveragePooling2D)	(None, 6, 6, 16)	0
conv2d_2 (Conv2D)	(None, 2, 2, 120)	48120
flatten (Flatten)	(None, 480)	0

```
In [10]: lenet.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 6)	456
average_pooling2d (AveragePooling2D)	(None, 16, 16, 6)	0
conv2d_1 (Conv2D)	(None, 12, 12, 16)	2416
average_pooling2d_1 (AveragePooling2D)	(None, 6, 6, 16)	0
conv2d_2 (Conv2D)	(None, 2, 2, 120)	48120
flatten (Flatten)	(None, 480)	0
dense (Dense)	(None, 84)	40404
dense_1 (Dense)	(None, 10)	850
Total params: 92,246		
Trainable params: 92,246		
Non-trainable params: 0		



```
In [11]: lenet.compile(optimizer='adam', loss=keras.losses.sparse_categorical_crossentropy, metrics=['accuracy'])
```

Neural Networks Deep Learning – Icp6

```
In [12]: hist = lenet.fit(x_train, y_train, epochs=100, validation_data=(x_test, y_test), verbose=1)
1563/1563 [=====] - 9s 6ms/step - loss: 0.0966 - accuracy: 0.9683 - val_loss: 4.9078 - val_accu
acy: 0.5907
Epoch 95/100
1563/1563 [=====] - 9s 6ms/step - loss: 0.0966 - accuracy: 0.9683 - val_loss: 4.9078 - val_accu
acy: 0.5959
Epoch 96/100
1563/1563 [=====] - 8s 5ms/step - loss: 0.1112 - accuracy: 0.9635 - val_loss: 4.8539 - val_accu
acy: 0.5863
Epoch 97/100
1563/1563 [=====] - 8s 5ms/step - loss: 0.1035 - accuracy: 0.9671 - val_loss: 5.0162 - val_accu
acy: 0.5881
Epoch 98/100
1563/1563 [=====] - 8s 5ms/step - loss: 0.1029 - accuracy: 0.9678 - val_loss: 4.9166 - val_accu
acy: 0.5819
Epoch 99/100
1563/1563 [=====] - 10s 6ms/step - loss: 0.1059 - accuracy: 0.9658 - val_loss: 5.0741 - val_accu
acy: 0.5924
Epoch 100/100
1563/1563 [=====] - 8s 5ms/step - loss: 0.1006 - accuracy: 0.9670 - val_loss: 4.9955 - val_accu
acy: 0.5902
```

```
In [13]: # summarize history for accuracy
plt.plot(hist.history['accuracy'])
plt.plot(hist.history['val_accuracy'])
plt.title("Accuracy by LeNet on CIFAR-10 Data")
plt.ylabel('Accuracy')
plt.xlabel('Epochs')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
# summarize history for loss
plt.plot(hist.history['loss'])
plt.plot(hist.history['val_loss'])
plt.title('Loss by LeNet on CIFAR-10 Data')
plt.ylabel('Loss')
plt.xlabel('Epochs')
plt.legend(['Train', 'Validation'])
plt.show()
```

```
plt.legend(['Train', 'Validation'])
plt.show()
```



Neural Networks Deep Learning – Icp6


```
In [14]: from sklearn.metrics import confusion_matrix
from sklearn.metrics import ConfusionMatrixDisplay
y_predictions= lenet.predict(x_test)
y_predictions.reshape(-1,)
y_predictions= np.argmax(y_predictions, axis=1)

confusion_matrix(y_test, y_predictions)
```

313/313 [=====] - 1s 2ms/step

```
Out[14]: array([[618, 36, 64, 32, 28, 21, 11, 24, 115, 51],
 [ 24, 740, 12, 17, 7, 10, 12, 8, 43, 127],
 [ 69, 21, 495, 76, 77, 110, 53, 55, 21, 23],
 [ 25, 27, 83, 397, 61, 215, 67, 63, 23, 39],
 [ 32, 10, 114, 90, 441, 80, 75, 125, 12, 21],
 [ 18, 8, 75, 200, 39, 523, 38, 70, 13, 16],
 [ 12, 23, 82, 102, 40, 76, 615, 13, 21, 16],
 [ 25, 7, 37, 65, 68, 95, 8, 664, 8, 23],
 [ 77, 55, 11, 22, 12, 16, 14, 9, 737, 47],
 [ 33, 142, 13, 30, 11, 23, 9, 29, 38, 672]])
```

```
In [16]: # confusion matrix and accuracy
from sklearn.metrics import confusion_matrix, accuracy_score
plt.figure(figsize=(7, 6))
plt.title('Confusion matrix', fontsize=16)
plt.imshow(confusion_matrix(y_test, y_predictions))
plt.xticks(np.arange(10), classes, rotation=45, fontsize=12)
plt.yticks(np.arange(10), classes, fontsize=12)
plt.colorbar()
plt.show()
```



```
print("Test accuracy:", accuracy_score(y_test, y_predictions))
```

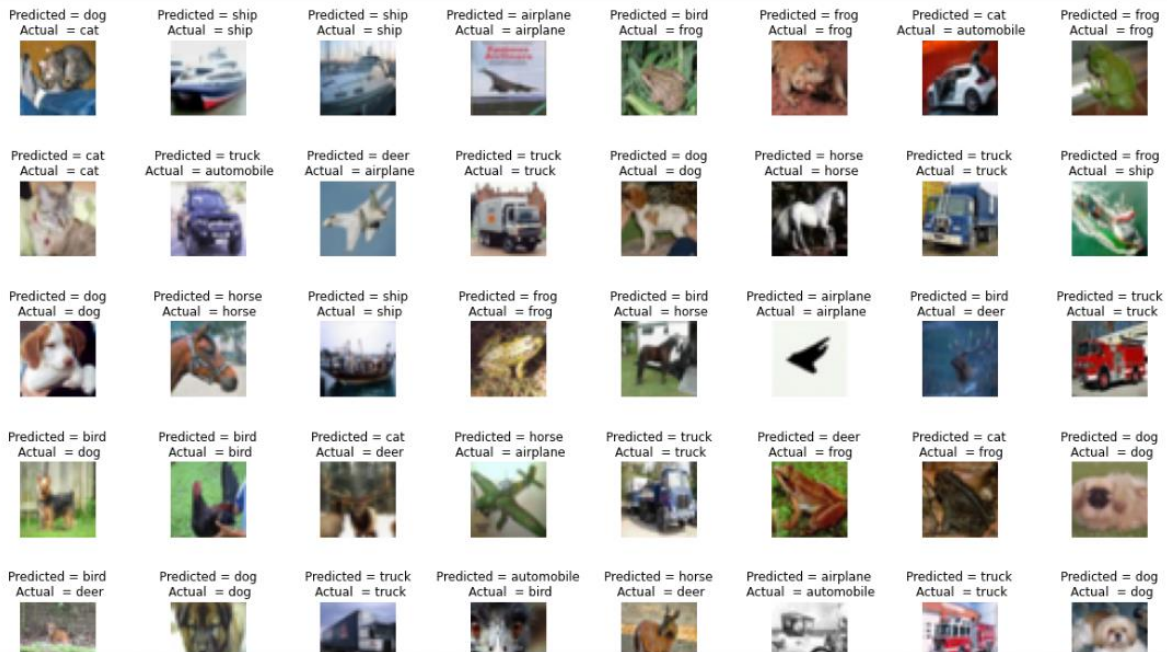
Test accuracy: 0.5902

Neural Networks Deep Learning – Icp6

```
19]: L = 8
W = 8
fig, axes = plt.subplots(L, W, figsize = (20,20))
axes = axes.ravel() #

for i in np.arange(0, L * W):
    axes[i].imshow(x_test[i])
    axes[i].set_title("Predicted = {}\n Actual = {}".format(classes[y_predictions[i]], classes[y_test[i]]))
    axes[i].axis('off')

plt.subplots_adjust(wspace=1)
```



```
: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Conv2D
from tensorflow.keras.layers import Dropout
from tensorflow.keras.layers import Flatten
from tensorflow.keras.optimizers import SGD, Adam
from keras.layers.convolutional import Convolution2D
from keras.layers.convolutional import MaxPooling2D

#Define Alexnet Model
AlexNet = Sequential()
AlexNet.add(Conv2D(filters=16,kernel_size=(3,3),strides=(4,4),input_shape=(32,32,3), activation='relu'))
AlexNet.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
AlexNet.add(Conv2D(60,(5,5),padding='same',activation='relu'))
AlexNet.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
AlexNet.add(Conv2D(60,(3,3),padding='same',activation='relu'))
AlexNet.add(Conv2D(30,(3,3),padding='same',activation='relu'))
AlexNet.add(Conv2D(20,(3,3),padding='same',activation='relu'))
AlexNet.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
AlexNet.add(Flatten())
AlexNet.add(Dense(200, activation='relu'))
AlexNet.add(Dropout(0.1))
AlexNet.add(Dense(200, activation='relu'))
AlexNet.add(Dropout(0.1))
AlexNet.add(Dense(10,activation='softmax'))

AlexNet.compile(optimizer='SGD', loss=keras.losses.sparse_categorical_crossentropy, metrics=['accuracy'])
AlexNet.summary()
```

Neural Networks Deep Learning – lcp6

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 8, 8, 16)	448
max_pooling2d (MaxPooling2D)	(None, 4, 4, 16)	0
conv2d_4 (Conv2D)	(None, 4, 4, 60)	24060
max_pooling2d_1 (MaxPooling2D)	(None, 2, 2, 60)	0
conv2d_5 (Conv2D)	(None, 2, 2, 60)	32460
conv2d_6 (Conv2D)	(None, 2, 2, 30)	16230
conv2d_7 (Conv2D)	(None, 2, 2, 20)	5420
max_pooling2d_2 (MaxPooling2D)	(None, 1, 1, 20)	0
flatten_1 (Flatten)	(None, 20)	0
dense_2 (Dense)	(None, 200)	4200
dropout (Dropout)	(None, 200)	0
dense_3 (Dense)	(None, 200)	40200
dropout_1 (Dropout)	(None, 200)	0
dense_4 (Dense)	(None, 10)	2010

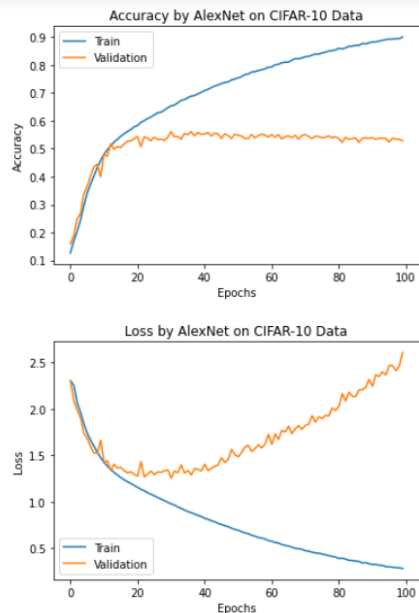
=====
Total params: 125,028
Trainable params: 125,028
Non-trainable params: 0
=====

```
[22]: history1 = AlexNet.fit(x_train, y_train, epochs=100, validation_data=(x_test, y_test), verbose=1)
racy: 0.5372
Epoch 95/100
1563/1563 [=====] - 10s 6ms/step - loss: 0.3035 - accuracy: 0.8912 - val_loss: 2.3625 - val_accu
racy: 0.5365
Epoch 96/100
1563/1563 [=====] - 10s 6ms/step - loss: 0.3017 - accuracy: 0.8921 - val_loss: 2.4652 - val_accu
racy: 0.5241
Epoch 97/100
1563/1563 [=====] - 11s 7ms/step - loss: 0.2984 - accuracy: 0.8927 - val_loss: 2.4627 - val_accu
racy: 0.5373
Epoch 98/100
1563/1563 [=====] - 11s 7ms/step - loss: 0.2925 - accuracy: 0.8944 - val_loss: 2.4082 - val_accu
racy: 0.5349
Epoch 99/100
1563/1563 [=====] - 10s 7ms/step - loss: 0.2918 - accuracy: 0.8943 - val_loss: 2.4604 - val_accu
racy: 0.5338
Epoch 100/100
1563/1563 [=====] - 11s 7ms/step - loss: 0.2835 - accuracy: 0.9001 - val_loss: 2.6043 - val_accu
racy: 0.5287
```

```
[23]: # summarize history for accuracy
plt.plot(history1.history['accuracy'])
plt.plot(history1.history['val_accuracy'])
plt.title("Accuracy by AlexNet on CIFAR-10 Data")
plt.ylabel('Accuracy')
plt.xlabel('Epochs')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()


# summarize history for loss
plt.plot(history1.history['loss'])
plt.plot(history1.history['val_loss'])
plt.title('Loss by AlexNet on CIFAR-10 Data')
plt.ylabel('Loss')
plt.xlabel('Epochs')
plt.legend(['Train', 'Validation'])
plt.show()
```

Neural Networks Deep Learning – Icp6



```
[24]: y_predictions1 = AlexNet.predict(x_test)
      y_predictions1.reshape(-1,)
      y_predictions1 = np.argmax(y_predictions1, axis=1)
      confusion_matrix(y_test, y_predictions1)

313/313 [=====] - 1s 2ms/step
```



```
[26]: print("Test accuracy by AlexNet:", accuracy_score(y_test, y_predictions))

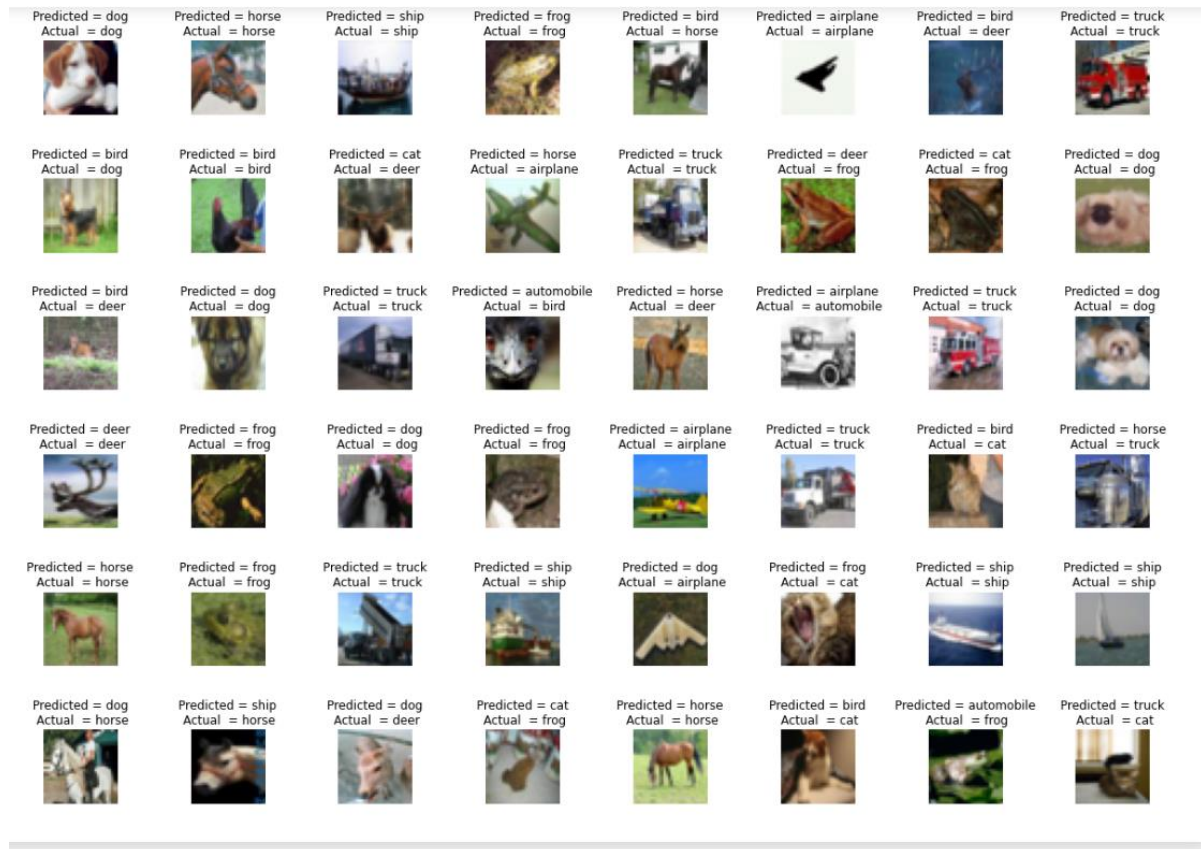
Test accuracy by AlexNet: 0.5902
```

```
[27]: L = 8
      W = 8
      fig, axes = plt.subplots(L, W, figsize = (20,20))
      axes = axes.ravel() #

      for i in np.arange(0, L * W):
          axes[i].imshow(x_test[i])
          axes[i].set_title("Predicted = {}\n Actual = {}".format(classes[y_predictions[i]], classes[y_test[i]]))
          axes[i].axis('off')

      plt.subplots_adjust(wspace=1)
```


Neural Networks Deep Learning – Icp6



VGG16

```
import keras
from keras.models import Sequential
from keras.preprocessing import image
from keras.layers import Activation,Dense,Dropout,Conv2D,Flatten,MaxPooling2D,BatchNormalization
from keras.datasets import cifar10
from keras import optimizers
from matplotlib import pyplot as plt
```

```
# generate cifar10 data
(x_train,y_train),(x_test,y_test) = cifar10.load_data()
```

```
# config parameters
num_classes = 10
input_shape = x_train.shape[1:4]
optimizer = optimizers.Adam(lr=0.0003)
```

WARNING:absl:lr is deprecated in Keras optimizer, please use learning_rate or use the legacy optimizer, e.g.,tf.keras.optimizers.legacy.Adam.

```
# convert label to one-hot
one_hot_y_train = keras.utils.to_categorical(y_train,num_classes=num_classes)
one_hot_y_test = keras.utils.to_categorical(y_test,num_classes=num_classes)
```


```
# check data
plt.imshow(x_train[1])
print(x_train[1].shape)
```


Neural Networks Deep Learning – Icp6

```
one_hot_y_test = keras.utils.to_categorical(y_test,num_classes=num_classes)
```

```
22]: # check data
plt.imshow(x_train[1])
print(x_train[1].shape)
```

(32, 32, 3)



```
# build model(similar to VGG16, only change the input and output shape)
model = Sequential()
model.add(Conv2D(64,(3,3),activation='relu',input_shape=input_shape,padding='same'))
model.add(Conv2D(64,(3,3),activation='relu',padding='same'))
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Conv2D(128,(3,3),activation='relu',padding='same'))
model.add(Conv2D(128,(3,3),activation='relu',padding='same'))
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Conv2D(256,(3,3),activation='relu',padding='same'))
model.add(Conv2D(256,(3,3),activation='relu',padding='same'))
model.add(Conv2D(256,(3,3),activation='relu',padding='same'))
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Flatten())
model.add(Dense(4096,activation='relu'))
model.add(Dense(4096,activation='relu'))
model.add(Dense(num_classes))
model.add(Activation('softmax'))
```

```
# config optimizer,loss,metrics
model.compile(optimizer=optimizer,loss='categorical_crossentropy',metrics=['accuracy'])
```

Neural Networks Deep Learning – lcp6

```
27]: # train
      history=model.fit(x=x_train,y=one_hot_y_train,batch_size=128,epochs=10,validation_split=0.1)

Epoch 1/10
352/352 [=====] - 25s 70ms/step - loss: 2.3027 - accuracy: 0.0997 - val_loss: 2.3027 - val_accuracy:
0.0950
Epoch 2/10
352/352 [=====] - 22s 61ms/step - loss: 2.3027 - accuracy: 0.0993 - val_loss: 2.3027 - val_accuracy:
0.0986
Epoch 3/10
352/352 [=====] - 21s 61ms/step - loss: 2.3027 - accuracy: 0.0998 - val_loss: 2.3028 - val_accuracy:
0.0950
Epoch 4/10
352/352 [=====] - 22s 62ms/step - loss: 2.3027 - accuracy: 0.0993 - val_loss: 2.3029 - val_accuracy:
0.0950
Epoch 5/10
352/352 [=====] - 22s 63ms/step - loss: 2.3027 - accuracy: 0.0965 - val_loss: 2.3029 - val_accuracy:
0.0958
Epoch 6/10
352/352 [=====] - 22s 62ms/step - loss: 2.3027 - accuracy: 0.0955 - val_loss: 2.3028 - val_accuracy:
0.0976
Epoch 7/10
352/352 [=====] - 22s 61ms/step - loss: 2.3027 - accuracy: 0.0990 - val_loss: 2.3029 - val_accuracy:
0.0958
Epoch 8/10
352/352 [=====] - 21s 61ms/step - loss: 2.3027 - accuracy: 0.0983 - val_loss: 2.3028 - val_accuracy:
0.1024
Epoch 9/10
352/352 [=====] - 21s 61ms/step - loss: 2.3027 - accuracy: 0.0990 - val_loss: 2.3027 - val_accuracy:
0.0976
Epoch 10/10
352/352 [=====] - 21s 61ms/step - loss: 2.3027 - accuracy: 0.0990 - val_loss: 2.3028 - val_accuracy:
0.0976

281:
```

Neural Networks Deep Learning – Icp6


```
]: # predict
plt.imshow(x_test[1000])

result = model.predict(x_test[1000:1001]).tolist()
predict = 0
expect = y_test[1000][0]
for i, _ in enumerate(result[0]):
    if result[0][i] > result[0][predict]:
        predict = i
print("predict class:", predict)
print("expected class:", expect)
```


1/1 [=====] - 0s 143ms/step

predict class: 6

expected class: 5




```
plt.plot(history.history['accuracy'], label='training accuracy')
plt.plot(history.history['val_accuracy'], label='validation accuracy')
plt.title('Training and validation accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```



Neural Networks Deep Learning – Icp6

```
[32]: plt.plot(history.history['loss'], label='training loss')
plt.plot(history.history['val_loss'], label='validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.show()
```




```
import numpy as np
from sklearn.metrics import confusion_matrix

# calculate the confusion matrix
y_pred = model.predict(x_test)
y_pred_classes = np.argmax(y_pred, axis=1)
y_true = y_test.ravel()
cm = confusion_matrix(y_true, y_pred_classes)


# plot the confusion matrix
plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
plt.title('Confusion matrix')
plt.colorbar()
tick_marks = np.arange(num_classes)
plt.xticks(tick_marks, range(num_classes))
plt.yticks(tick_marks, range(num_classes))
plt.xlabel('Predicted label')
plt.ylabel('True label')
plt.show()

# plot a histogram of the predicted probabilities for a sample image
plt.hist(y_pred[1000])
plt.title('Predicted probabilities')
plt.xlabel('Probability')
plt.ylabel('Frequency')
plt.show()
```

```
313/313 [=====] - 4s 9ms/step
```



Neural Networks Deep Learning – lcp6



VGG19MODE_CIFAR100:

Import libraries

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import tensorflow as tf
```

In []:

```
from tensorflow.keras.optimizers import RMSprop
from keras.preprocessing import image
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D, Dropout, BatchNormalization
```

```
%matplotlib inline
```

Extract data and train and test dataset

```
cifar100 = tf.keras.datasets.cifar100
(X_train,Y_train) , (X_test,Y_test) = cifar10.load_data()
```

In []:

```
classes = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
```


In []:

Let's look into the dataset images

In []:

```
plt.figure(figsize = (16,16))
for i in range(100):
    plt.subplot(10,10,1+i)
    plt.axis('off')
    plt.imshow(X_train[i], cmap = 'gray')
```

Neural Networks Deep Learning – Icp6



Training , Validating and Splitting trained and tested data

```
from sklearn.model_selection import train_test_split
x_train, x_val, y_train, y_val = train_test_split(X_train, Y_train, test_size=0.2)
```

```
from keras.utils.np_utils import to_categorical
y_train = to_categorical(y_train, num_classes = 10)
y_val = to_categorical(y_val, num_classes = 10)
```

```
print(x_train.shape)
print(y_train.shape)
```

In []:

In []:

In []:

Neural Networks Deep Learning – lcp6

```
print(x_val.shape)
print(y_val.shape)
print(X_test.shape)
print(Y_test.shape)
```

```
(40000, 32, 32, 3)
(40000, 10)
(10000, 32, 32, 3)
(10000, 10)
(10000, 32, 32, 3)
(10000, 1)
```

In []:

```
train_datagen = ImageDataGenerator(
    preprocessing_function = tf.keras.applications.vgg19.preprocess_input,
    rotation_range=10,
    zoom_range = 0.1,
    width_shift_range = 0.1,
    height_shift_range = 0.1,
    shear_range = 0.1,
    horizontal_flip = True
)
train_datagen.fit(x_train)
```

```
val_datagen = ImageDataGenerator(preprocessing_function = tf.keras.applications.vgg19.preprocess_input)
val_datagen.fit(x_val)
```

In []:

```
from keras.callbacks import ReduceLROnPlateau
learning_rate_reduction = ReduceLROnPlateau(monitor='val_accuracy',
                                             patience=3,
                                             verbose=1,
                                             factor=0.5,
                                             min_lr=0.00001)
```

We have used only 16 layers out of 19 layers in the CNN

In []:

```
vgg_model = tf.keras.applications.VGG19(
    include_top=False,
    weights="imagenet",
    input_shape=(32,32,3),
)
```

```
vgg_model.summary()
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg19/vgg19_weights_tf_dim_ordering_tf_kernels_notop.h5

80142336/80134624 [=====] - 1s 0us/step

Model: "vgg19"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 32, 32, 3)]	0

block1_conv1 (Conv2D)	(None, 32, 32, 64)	1792

block1_conv2 (Conv2D)	(None, 32, 32, 64)	36928

Neural Networks Deep Learning – lcp6

block1_pool (MaxPooling2D)	(None, 16, 16, 64)	0
block2_conv1 (Conv2D)	(None, 16, 16, 128)	73856
block2_conv2 (Conv2D)	(None, 16, 16, 128)	147584
block2_pool (MaxPooling2D)	(None, 8, 8, 128)	0
block3_conv1 (Conv2D)	(None, 8, 8, 256)	295168
block3_conv2 (Conv2D)	(None, 8, 8, 256)	590080
block3_conv3 (Conv2D)	(None, 8, 8, 256)	590080
block3_conv4 (Conv2D)	(None, 8, 8, 256)	590080
block3_pool (MaxPooling2D)	(None, 4, 4, 256)	0
block4_conv1 (Conv2D)	(None, 4, 4, 512)	1180160
block4_conv2 (Conv2D)	(None, 4, 4, 512)	2359808
block4_conv3 (Conv2D)	(None, 4, 4, 512)	2359808
block4_conv4 (Conv2D)	(None, 4, 4, 512)	2359808
block4_pool (MaxPooling2D)	(None, 2, 2, 512)	0
block5_conv1 (Conv2D)	(None, 2, 2, 512)	2359808
block5_conv2 (Conv2D)	(None, 2, 2, 512)	2359808
block5_conv3 (Conv2D)	(None, 2, 2, 512)	2359808
block5_conv4 (Conv2D)	(None, 2, 2, 512)	2359808
block5_pool (MaxPooling2D)	(None, 1, 1, 512)	0
=====		
Total params: 20,024,384		
Trainable params: 20,024,384		
Non-trainable params: 0		

```
model = tf.keras.Sequential()
model.add(vgg_model)
model.add(Flatten())
model.add(Dense(1024, activation = 'relu'))
model.add(Dense(1024, activation = 'relu'))
model.add(Dense(256, activation = 'relu'))
model.add(Dense(10, activation = 'softmax'))
```

```
model.summary()

Model: "sequential"
```

Layer (type)	Output Shape	Param #
--------------	--------------	---------

In []:

Neural Networks Deep Learning – lcp6

```
=====
vgg19 (Functional)      (None, 1, 1, 512)      20024384
```

```
flatten (Flatten)       (None, 512)            0
```

```
dense (Dense)           (None, 1024)           525312
```

```
dense_1 (Dense)         (None, 1024)           1049600
```

```
dense_2 (Dense)         (None, 256)            262400
```

```
dense_3 (Dense)         (None, 10)             2570
=====
```

Total params: 21,864,266

Trainable params: 21,864,266

Non-trainable params: 0

In []:

```
optimizer = tf.keras.optimizers.SGD(lr = 0.001, momentum = 0.9)
```

```
model.compile(optimizer= optimizer,
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

```
/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/optimizer_v2/optimizer_v2.py:375: UserWarning: The `lr` argument is deprecated, use `learning_rate` instead.
```

```
"The `lr` argument is deprecated, use `learning_rate` instead.")
```

In []:

```
history = model.fit(
    train_datagen.flow(x_train, y_train, batch_size = 128),
    validation_data = val_datagen.flow(x_val, y_val, batch_size = 128),
    epochs = 25,
    verbose = 1,
    callbacks = [learning_rate_reduction]
)
```

Epoch 1/25

313/313 [=====] - 84s 165ms/step - loss: 1.6671 - accuracy: 0.3875 - val_loss: 1.1241 - val_accuracy: 0.6056

Epoch 2/25

313/313 [=====] - 50s 159ms/step - loss: 0.9631 - accuracy: 0.6640 - val_loss: 0.7466 - val_accuracy: 0.7404

Epoch 3/25

313/313 [=====] - 50s 160ms/step - loss: 0.7464 - accuracy: 0.7430 - val_loss: 0.6792 - val_accuracy: 0.7716

Epoch 4/25

313/313 [=====] - 50s 160ms/step - loss: 0.6533 - accuracy: 0.7782 - val_loss: 0.6814 - val_accuracy: 0.7829

Epoch 5/25

313/313 [=====] - 50s 159ms/step - loss: 0.5779 - accuracy: 0.8013 - val_loss: 0.5932 - val_accuracy: 0.8058

Epoch 6/25

313/313 [=====] - 50s 160ms/step - loss: 0.5369 - accuracy: 0.8136 - val_loss: 0.5455 - val_accuracy: 0.8157

Epoch 7/25

313/313 [=====] - 50s 160ms/step - loss: 0.4925 - accuracy: 0.8299 - val_loss: 0.5119 - val_accuracy: 0.8269

Epoch 8/25

Neural Networks Deep Learning – lcp6

313/313 [=====] - 50s 160ms/step - loss: 0.4660 - accuracy: 0.8398 - val_loss: 0.5036 - val_accu
acy: 0.8342
Epoch 9/25
313/313 [=====] - 50s 160ms/step - loss: 0.4315 - accuracy: 0.8523 - val_loss: 0.4470 - val_accu
acy: 0.8461
Epoch 10/25
313/313 [=====] - 50s 160ms/step - loss: 0.4110 - accuracy: 0.8569 - val_loss: 0.4712 - val_accu
acy: 0.8440
Epoch 11/25
313/313 [=====] - 50s 159ms/step - loss: 0.3797 - accuracy: 0.8673 - val_loss: 0.4721 - val_accu
acy: 0.8504
Epoch 12/25
313/313 [=====] - 50s 160ms/step - loss: 0.3641 - accuracy: 0.8728 - val_loss: 0.4315 - val_accu
acy: 0.8541
Epoch 13/25
313/313 [=====] - 50s 159ms/step - loss: 0.3399 - accuracy: 0.8826 - val_loss: 0.4618 - val_accu
acy: 0.8501
Epoch 14/25
313/313 [=====] - 50s 159ms/step - loss: 0.3259 - accuracy: 0.8860 - val_loss: 0.4929 - val_accu
acy: 0.8448
Epoch 15/25
313/313 [=====] - 50s 159ms/step - loss: 0.3130 - accuracy: 0.8925 - val_loss: 0.4325 - val_accu
acy: 0.8604
Epoch 16/25
313/313 [=====] - 50s 159ms/step - loss: 0.2969 - accuracy: 0.8955 - val_loss: 0.4943 - val_accu
acy: 0.8477
Epoch 17/25
313/313 [=====] - 50s 159ms/step - loss: 0.2805 - accuracy: 0.9008 - val_loss: 0.4315 - val_accu
acy: 0.8650
Epoch 18/25
313/313 [=====] - 50s 159ms/step - loss: 0.2638 - accuracy: 0.9089 - val_loss: 0.4128 - val_accu
acy: 0.8674
Epoch 19/25
313/313 [=====] - 50s 159ms/step - loss: 0.2528 - accuracy: 0.9108 - val_loss: 0.4517 - val_accu
acy: 0.8616
Epoch 20/25
313/313 [=====] - 50s 159ms/step - loss: 0.2420 - accuracy: 0.9147 - val_loss: 0.4273 - val_accu
acy: 0.8658
Epoch 21/25
313/313 [=====] - 50s 159ms/step - loss: 0.2311 - accuracy: 0.9158 - val_loss: 0.4323 - val_accu
acy: 0.8693
Epoch 22/25
313/313 [=====] - 50s 159ms/step - loss: 0.2142 - accuracy: 0.9244 - val_loss: 0.4716 - val_accu
acy: 0.8679
Epoch 23/25
313/313 [=====] - 50s 159ms/step - loss: 0.2091 - accuracy: 0.9272 - val_loss: 0.4171 - val_accu
acy: 0.8765
Epoch 24/25
313/313 [=====] - 50s 159ms/step - loss: 0.1964 - accuracy: 0.9304 - val_loss: 0.4461 - val_accu
acy: 0.8724
Epoch 25/25
313/313 [=====] - 50s 159ms/step - loss: 0.1879 - accuracy: 0.9337 - val_loss: 0.4827 - val_accu
acy: 0.8671

In []:

```
acc = history.history['accuracy']
```

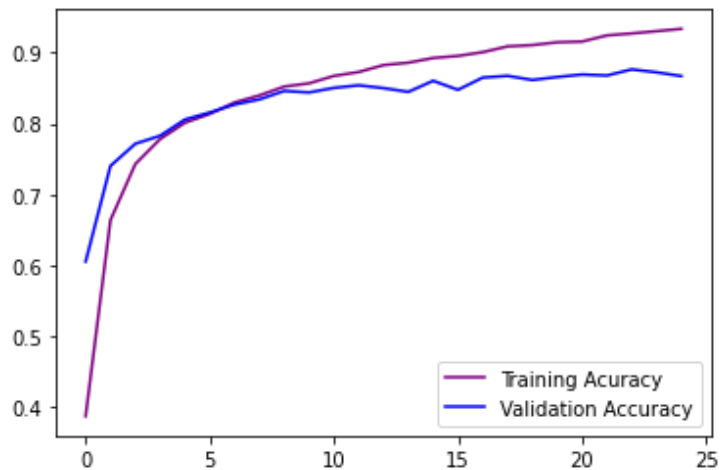
Neural Networks Deep Learning – lcp6

```
val_acc = history.history['val_accuracy']
```

```
plt.figure()  
plt.plot(acc,color = 'purple',label = 'Training Acuracy')  
plt.plot(val_acc,color = 'blue',label = 'Validation Accuracy')  
plt.legend()
```

Out []:

<matplotlib.legend.Legend at 0x7fc5601fe610>



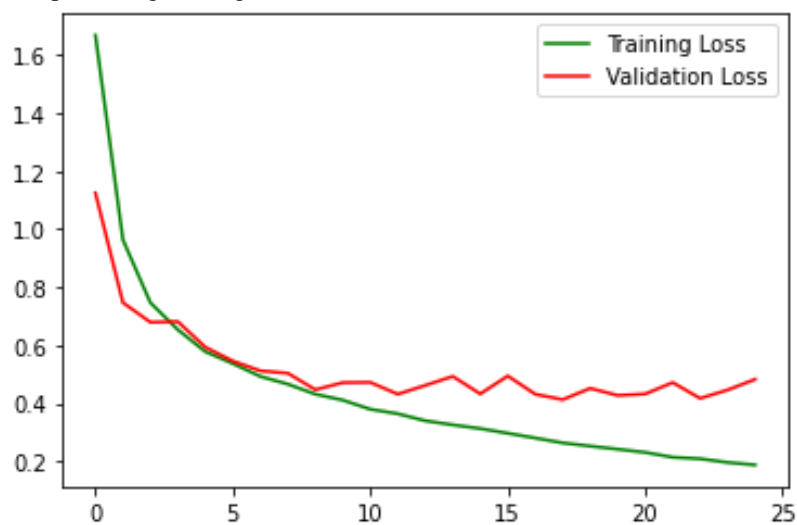
In []:

```
loss = history.history['loss']  
val_loss = history.history['val_loss']
```

```
plt.figure()  
plt.plot(loss,color = 'green',label = 'Training Loss')  
plt.plot(val_loss,color = 'red',label = 'Validation Loss')  
plt.legend()
```

Out []:

<matplotlib.legend.Legend at 0x7fc5bd878ad0>



In []:

```
x_test = tf.keras.applications.vgg19.preprocess_input(X_test)  
y_pred = model.predict_classes(x_test)  
y_pred[:10]
```

Neural Networks Deep Learning – Icp6

```
/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/engine/sequential.py:455: UserWarning: `model.predict_classes()` is deprecated and will be removed after 2021-01-01. Please use instead: * `np.argmax(model.predict(x), axis=-1)`, if your model does multi-class classification (e.g. if it uses a `softmax` last-layer activation). * `(model.predict(x) > 0.5).astype("int32")`, if your model does binary classification (e.g. if it uses a `sigmoid` last-layer activation).
```

```
warnings.warn("`model.predict_classes()` is deprecated and "
```

Out[]:

```
array([3, 8, 8, 0, 6, 6, 1, 6, 3, 1])
```

In []:

```
from sklearn.metrics import confusion_matrix, accuracy_score
print('Testing Accuracy : ', accuracy_score(Y_test, y_pred))
```

```
Testing Accuracy : 0.8591
```

In []:

```
cm = confusion_matrix(Y_test, y_pred)
cm
```

Out[]:

```
array([[882, 11, 20, 1, 12, 1, 5, 9, 43, 16],
       [ 5, 944, 0, 2, 0, 1, 2, 0, 8, 38],
       [22, 3, 800, 11, 54, 30, 56, 9, 10, 5],
       [10, 7, 42, 536, 57, 228, 74, 26, 6, 14],
       [ 6, 1, 22, 5, 884, 20, 35, 23, 2, 2],
       [ 5, 4, 9, 54, 36, 845, 19, 26, 0, 2],
       [ 5, 1, 11, 7, 16, 10, 941, 3, 3, 3],
       [ 4, 1, 12, 7, 31, 28, 8, 902, 2, 5],
       [25, 16, 1, 1, 4, 0, 4, 1, 929, 19],
       [ 6, 41, 1, 2, 4, 0, 6, 2, 10, 928]])
```

In []:

```
import itertools
```

```
def plot_confusion_matrix(cm, classes,
                           normalize=False,
                           title='Confusion matrix',
                           cmap=plt.cm.Greens):
    """
```

```
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
```

```
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=30)
    plt.yticks(tick_marks, classes)
```

```
    if normalize:
```

```
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
```

```
    else:
```

```
        print('Confusion matrix, without normalization')
```

```
    #print(cm)
```

```
    thresh = cm.max() / 2.
```

```
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
```

```
        plt.text(j, i, cm[i, j],
                  horizontalalignment="center",
                  color="white" if cm[i, j] > thresh else "black")
```

Neural Networks Deep Learning – Icp6

```
plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label')
```

```
plt.figure(figsize=(8,8))
plot_confusion_matrix(cm,classes)
```

Confusion matrix, without normalization

